

Image Segmentation

Introduction to Segment Anything

Jan Theiß

December 11, 2023

Outline

1 What are Transformers?

- Attention Is All You Need
- An Image Is Worth 16X16 Words

2 What is Segment Anything?

- Definition and Key Aspects
- What does SAM look like?
- SA-1B

3 Implementing SAM

- Hands-On Implementation

What are Transformers?

- We think about the concept of transformers most of the time in Natural Language Processing(NLP)-Tasks
- Transformers can modify (transform) the input
- A very basic transformer could be a translator

Attention Is All You Need

Target:

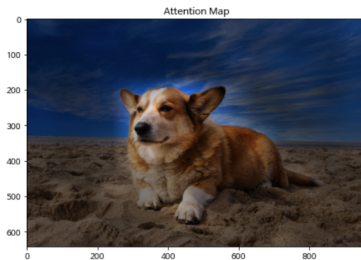
- Semantic Translation
- from word-for-word to sense-for-sense

⇒ pairwise inner product of every item of the input set (i.e. words in a sentence)

Problem when dealing with images as input:

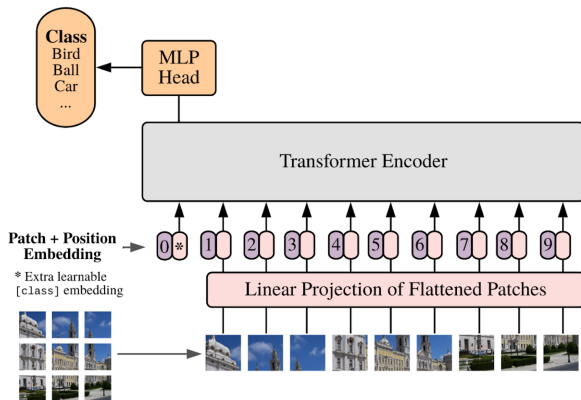
$250 \times 250 \Rightarrow 62.500 \text{ px} \Rightarrow 3.906.250.000 \text{ calculations}$

Attention Is All You Need

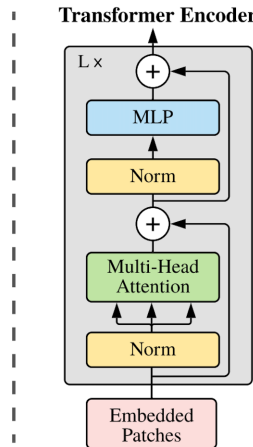


An Image Is Worth 16x16 Words

Vision Transformer (ViT)



Transformer Encoder



What is Segment Anything?

Video Demonstration

What is Segment Anything?

Definition And Key Aspects

Zero-Shot Large Language Model (LLM) developed by Meta AI that can segment objects in images

① Tasks

- ▶ How to interact with the model and what inputs are available

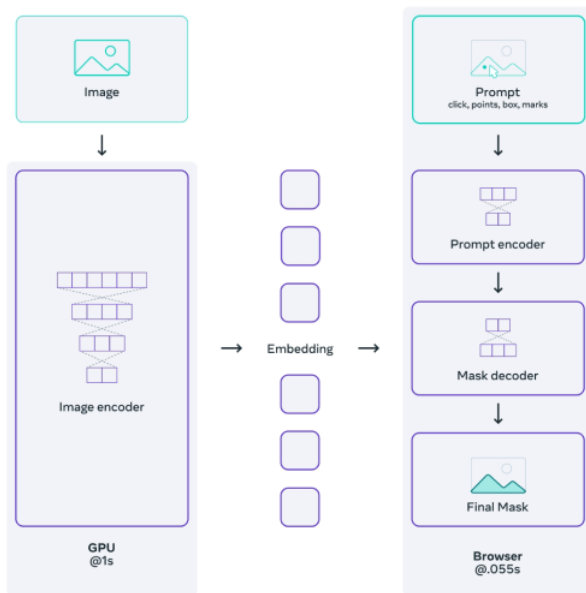
② Model

- ▶ Segment Anything Model (SAM)

③ Data

- ▶ SA-1B

What does SAM look like?



SA-1B Dataset

- 11M images
- 1.1B masks (99.1 % automatically collected)

DEMO

<https://segment-anything.com/demo>

Implementing SAM

- 1 go onto github
<https://github.com/facebookresearch/segment-anything/tree/main>
- 2 find an image of your choice
- 3 use opencv (or matplotlib.pyplot) to load the image
- 4 initiate SAM and a suitable predictor with the image
- 5 preview the masks